

TCP/IP Overview

Helpful References

- Forouzan, Data Communications and Networking, 5ed
- http://tcpipguide.com/free/t_toc.htm
- <https://www.geeksforgeeks.org/services-and-segment-structure-in-tcp/>

Network Review

- When it comes to designing, implementing, debugging, and testing networking software, we must be aware of
 - Network layers
 - Network headers
 - Network protocol flow

Network Layer Review

- Physical Layer:
 - Responsible for physical links between adjacent devices
 - Determines how to represent data in electromagnetic form
- Data Link Layer:
 - Responsible for getting data from one device to a neighbouring device in the current network
 - Breaks up data into individual frames
 - Handles error detection and correction as required
- Ethernet encompasses physical and data link layers

Network Layer Review

- Network Layer:
 - Responsible for determining network routes
 - Gets a packet from a source network to a destination network
 - IP
- Transport Layer:
 - Deals with sender and receiver issues: (De)multiplexing, packet ordering, flow control, etc.
 - TCP (reliable but slower), UDP (unreliable but faster)

IPv4 Header Review

Version	Header Length	DSCP	ECN	Header + Data Length	ID	Flags	Offset	TTL	Proto	Header Check sum	Src IP	Dst IP	Opts
---------	---------------	------	-----	----------------------	----	-------	--------	-----	-------	------------------	--------	--------	------

- Differentiated Services Code Point (DSCP) can be used to describe the payload (e.g., Voice over IP), allowing the router to make smarter choices
- Explicit Congestion Notification (ECN) can be used to alert the destination of impending congestion, which can then alert the source

IPv4 Header Review

Version	Header Length	DSCP	ECN	Header + Data Length	ID	Flags	Offset	TTL	Proto	Header Check sum	Src IP	Dst IP	Opts
---------	---------------	------	-----	----------------------	----	-------	--------	-----	-------	------------------	--------	--------	------

- ID + Flags + Offset used to enable/identify fragments
- Time To Live (TTL); usu. a per-router hop count; at 0, packet is discarded and an *ICMP Time Exceeded* error is transmitted to the sender
- Proto identifies the payload type

TCP Header Review

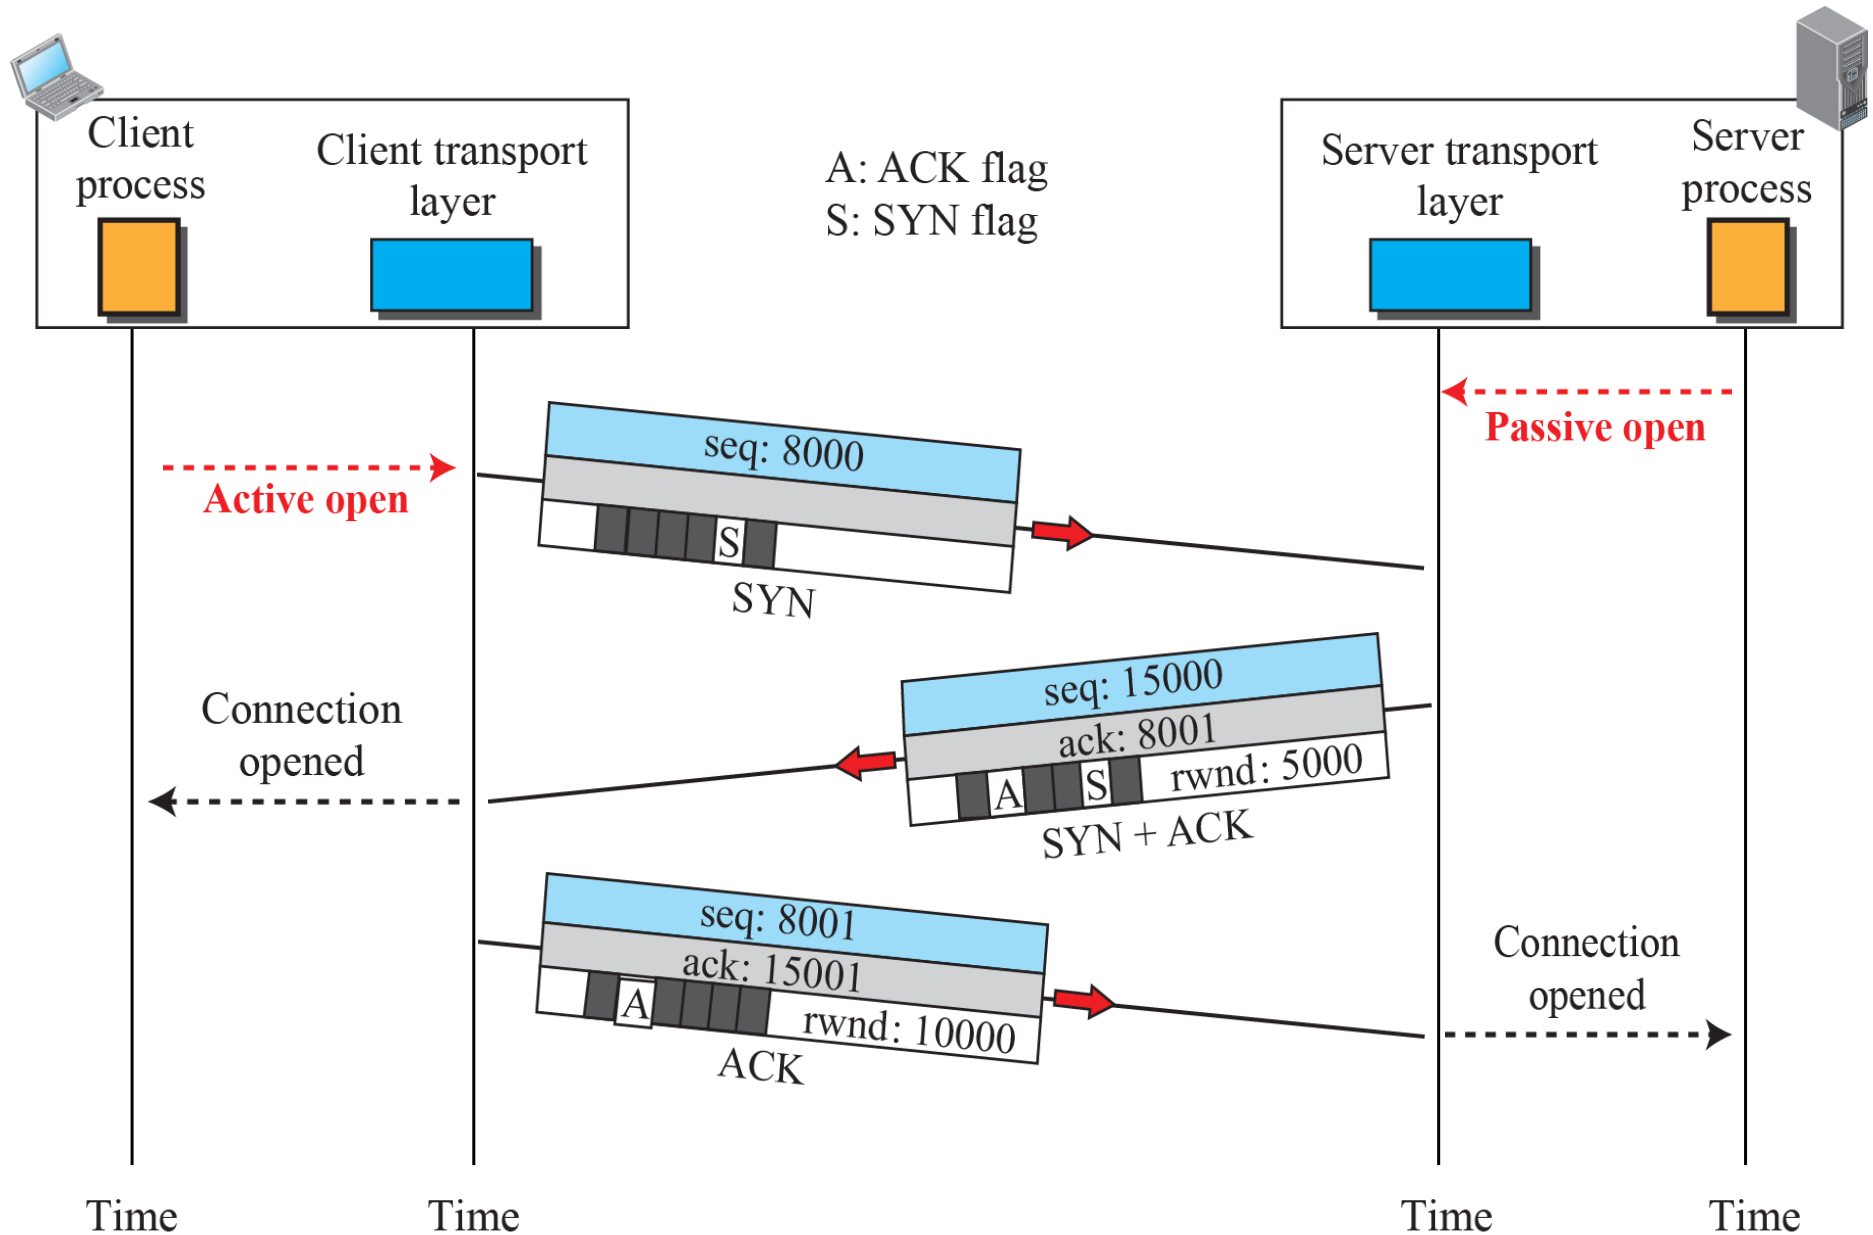
Src Port	Dst Port	Sequence Number	Ack. Number	Header Length	0	Flags	Window Size	Chksum	Urgent Pointer	Opts
----------	----------	-----------------	-------------	---------------	---	-------	-------------	--------	----------------	------

- Urgent Pointer points to the location of the urgent data in the payload section (if URG flag set)

UDP Header Review

Source Port	Destination Port	Header + Data Length	Checksum (Optional if used with IPv4)
-------------	------------------	----------------------	---------------------------------------

TCP Connection (3-way Handshake) Review



Connection Establishment Notes

- ACK = Acknowledgement
- rwnd = Receive window (used for flow control)
- SYN = Synchronize sequence numbers
- SYN and SYN/ACK segments do not carry data
- An ACK without data does not consume a sequence number
- Full duplex requires two receive windows

Sequence Numbers

- Initial number is random (Wireshark obscures this by default to make analysis easier)
- Subsequent numbers depend on number of bytes transmitted
 - e.g., assume first number is 10001
 - If 1000 bytes are transmitted, the sequence number range will be 10001 to 11000
 - The next segment would start at 11001

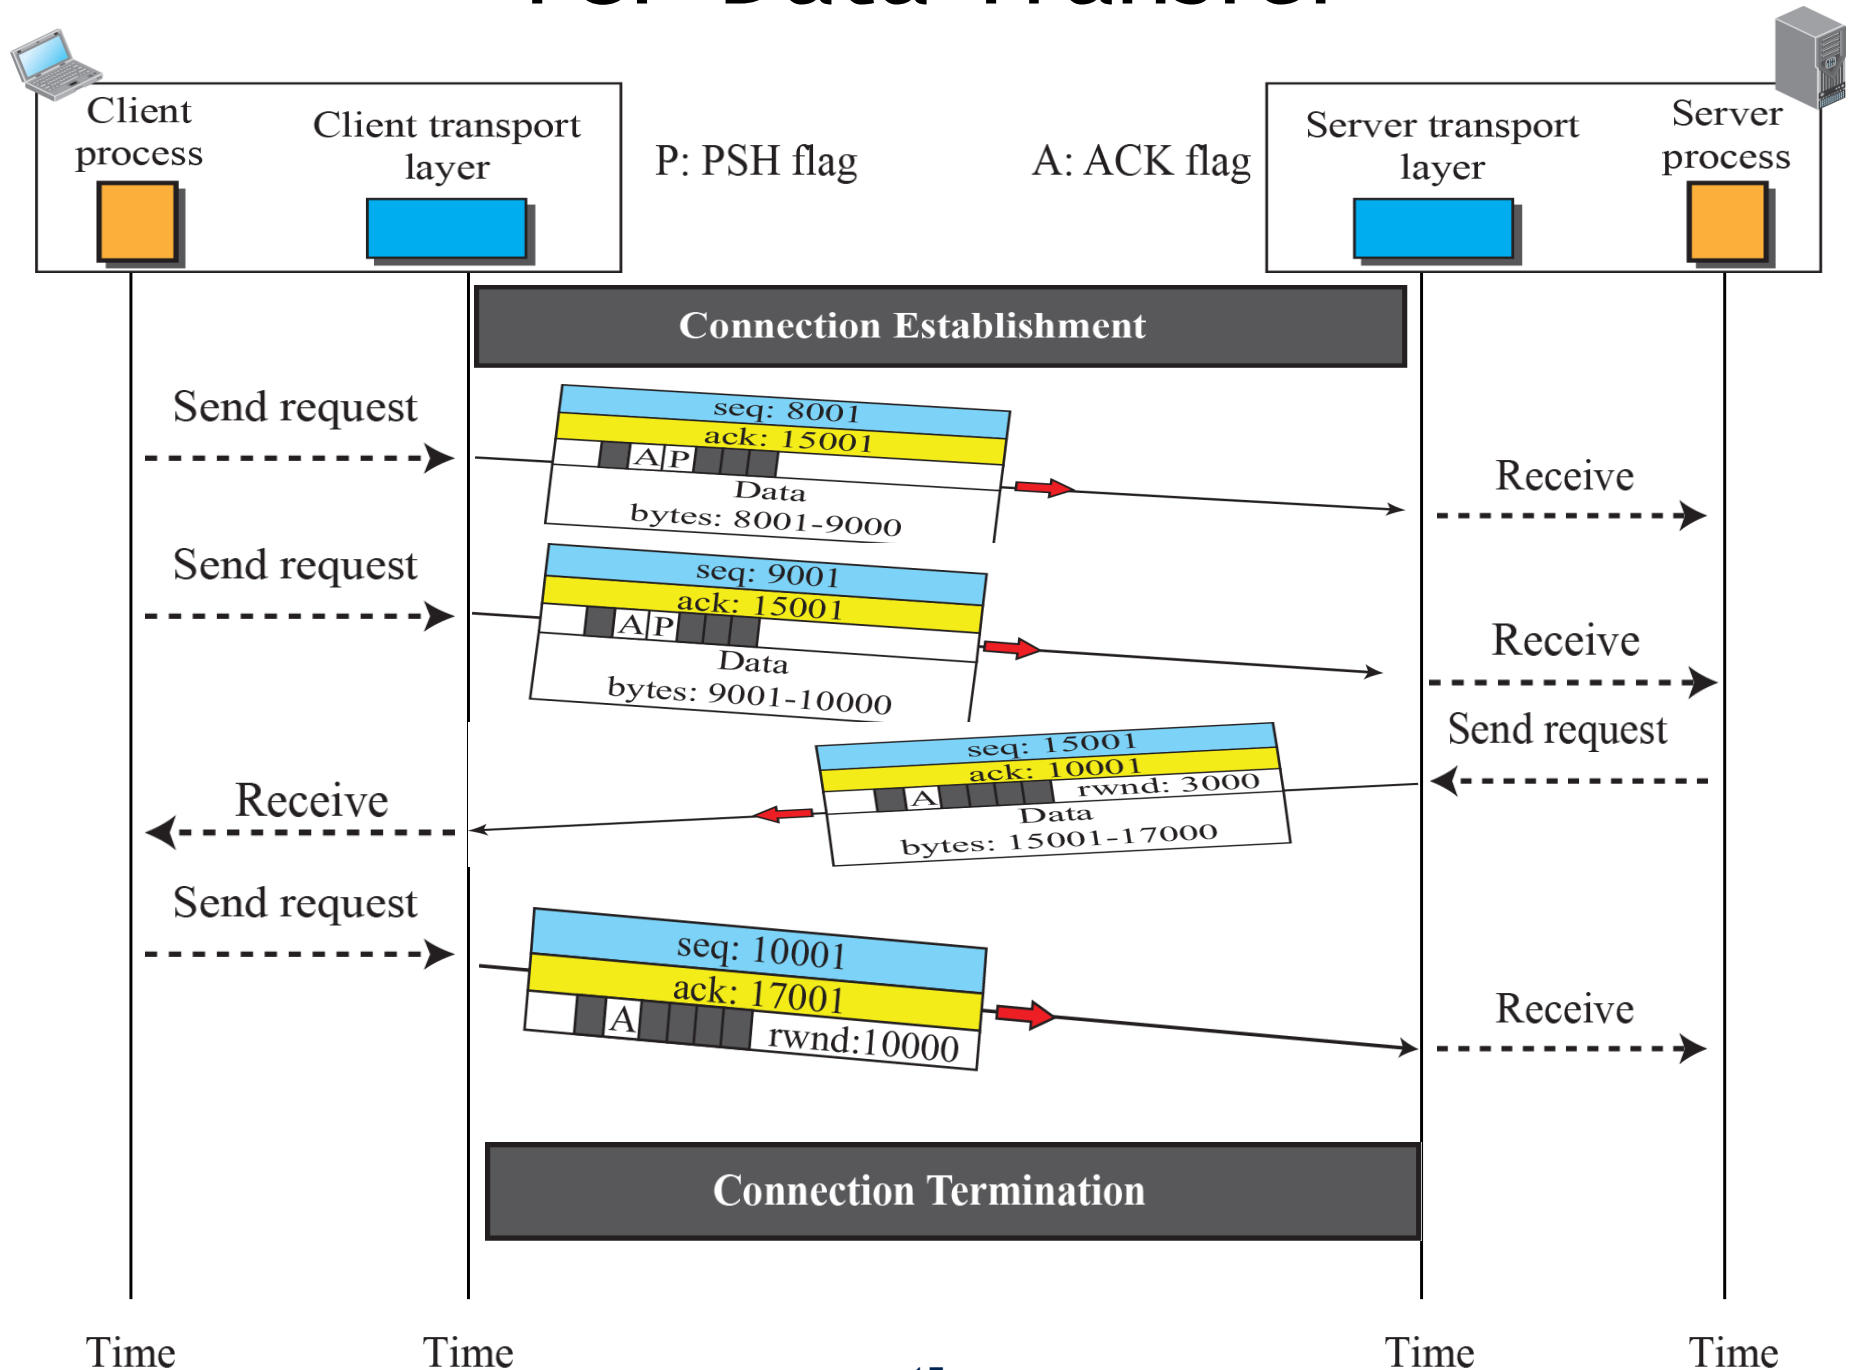
Exercises

- A TCP connection is established using a _____

Exercises

- During TCP connection establishment, if the server receives a packet with the SYN flag set and a sequence number set to 18573, it must reply with the _____ flag(s) set and the sequence number set to _____ and the acknowledgement number set to _____

TCP Data Transfer



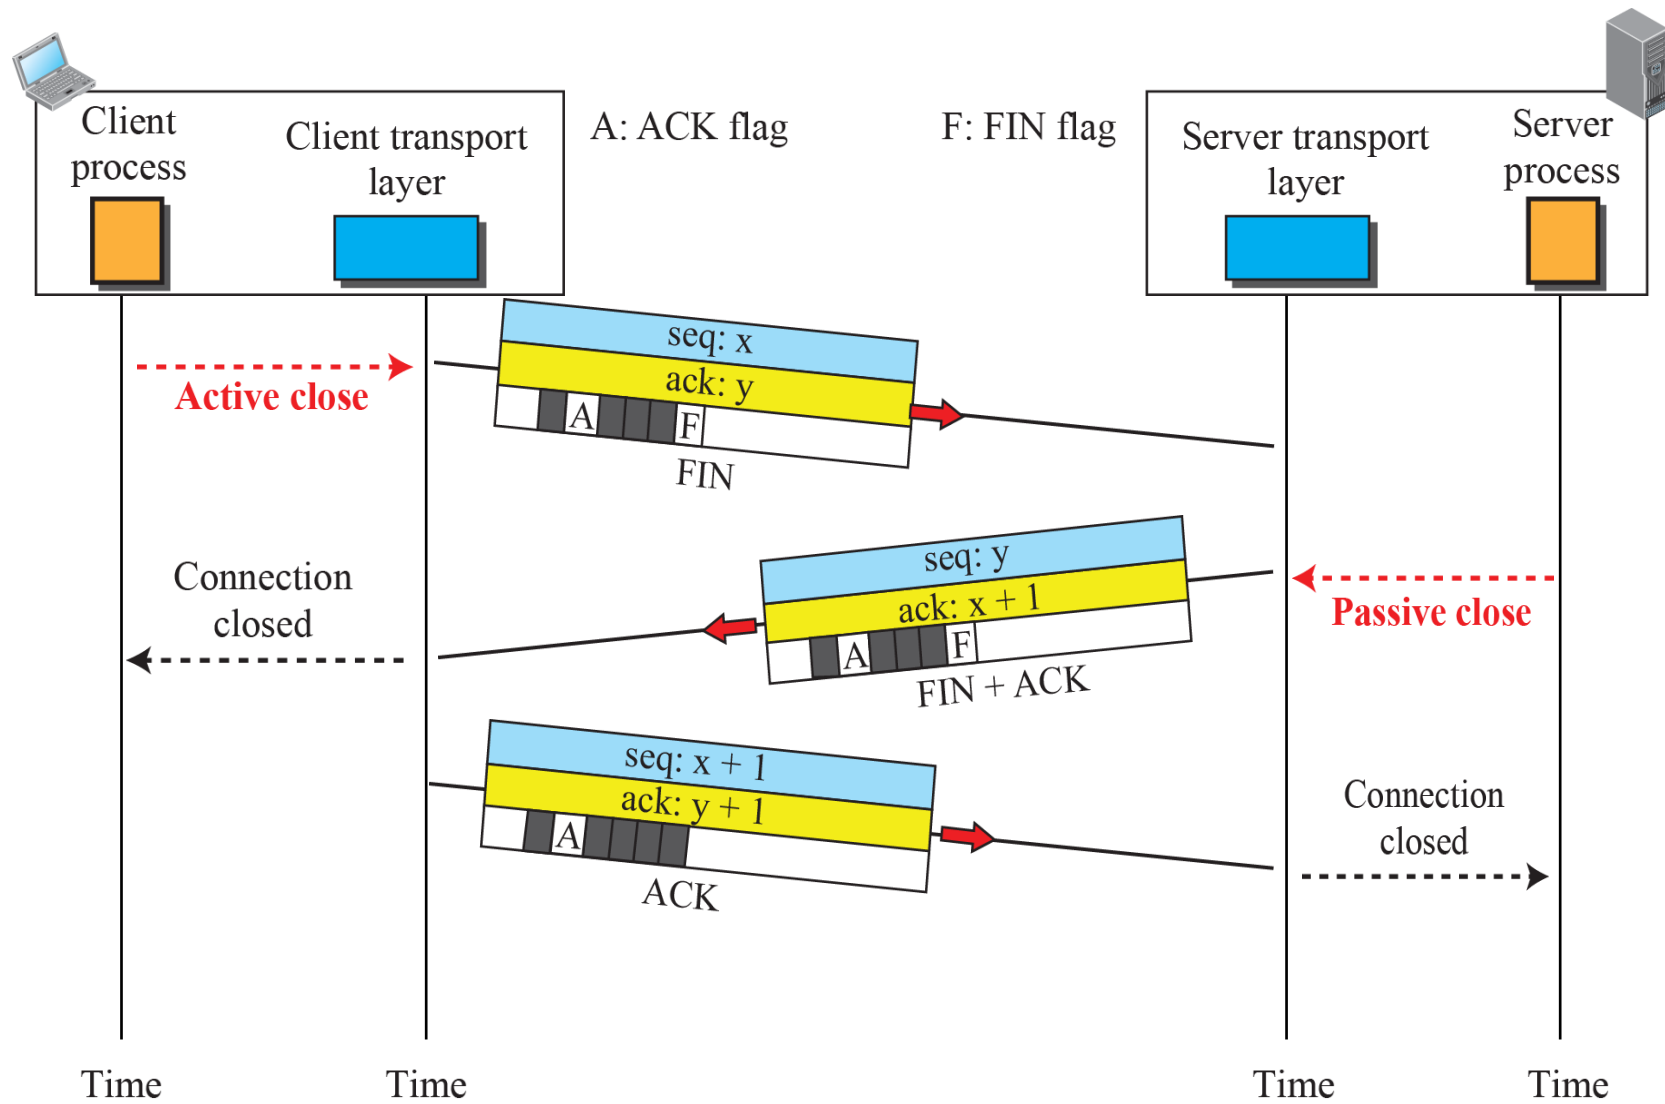
Data Transfer Notes

- The sequence number is referring to the first byte in the segment
- Note the window size update and the ACK number
- The push flag indicates to the destination that buffered data should be made available to the application right away, rather than delay for more data to accumulate
- Can emphasize urgency even more by setting the URG flag; points to urgent data in the data stream and causes a signal to be sent to the listening application

Exercises

- If the TCP segment sequence number is 58293 and 100 bytes are sent, the acknowledgement number must be _____


TCP Connection Termination



Connection Termination Notes


- FIN/FIN+ACK consumes 1 sequence number even if it contains no data
- Last ACK contains no data
- But what if more data must be sent? (e.g., a database query takes time to retrieve?)

TCP Half Close



Two Army Problem

- Assume the final ACK gets lost
- How do we know that the left hand side won't wait forever for a FIN?
- Problem: Do we ACK the ACK?



Source: Andy Tanenbaum, Computer Networks, 4th ed., Prentice-Hall

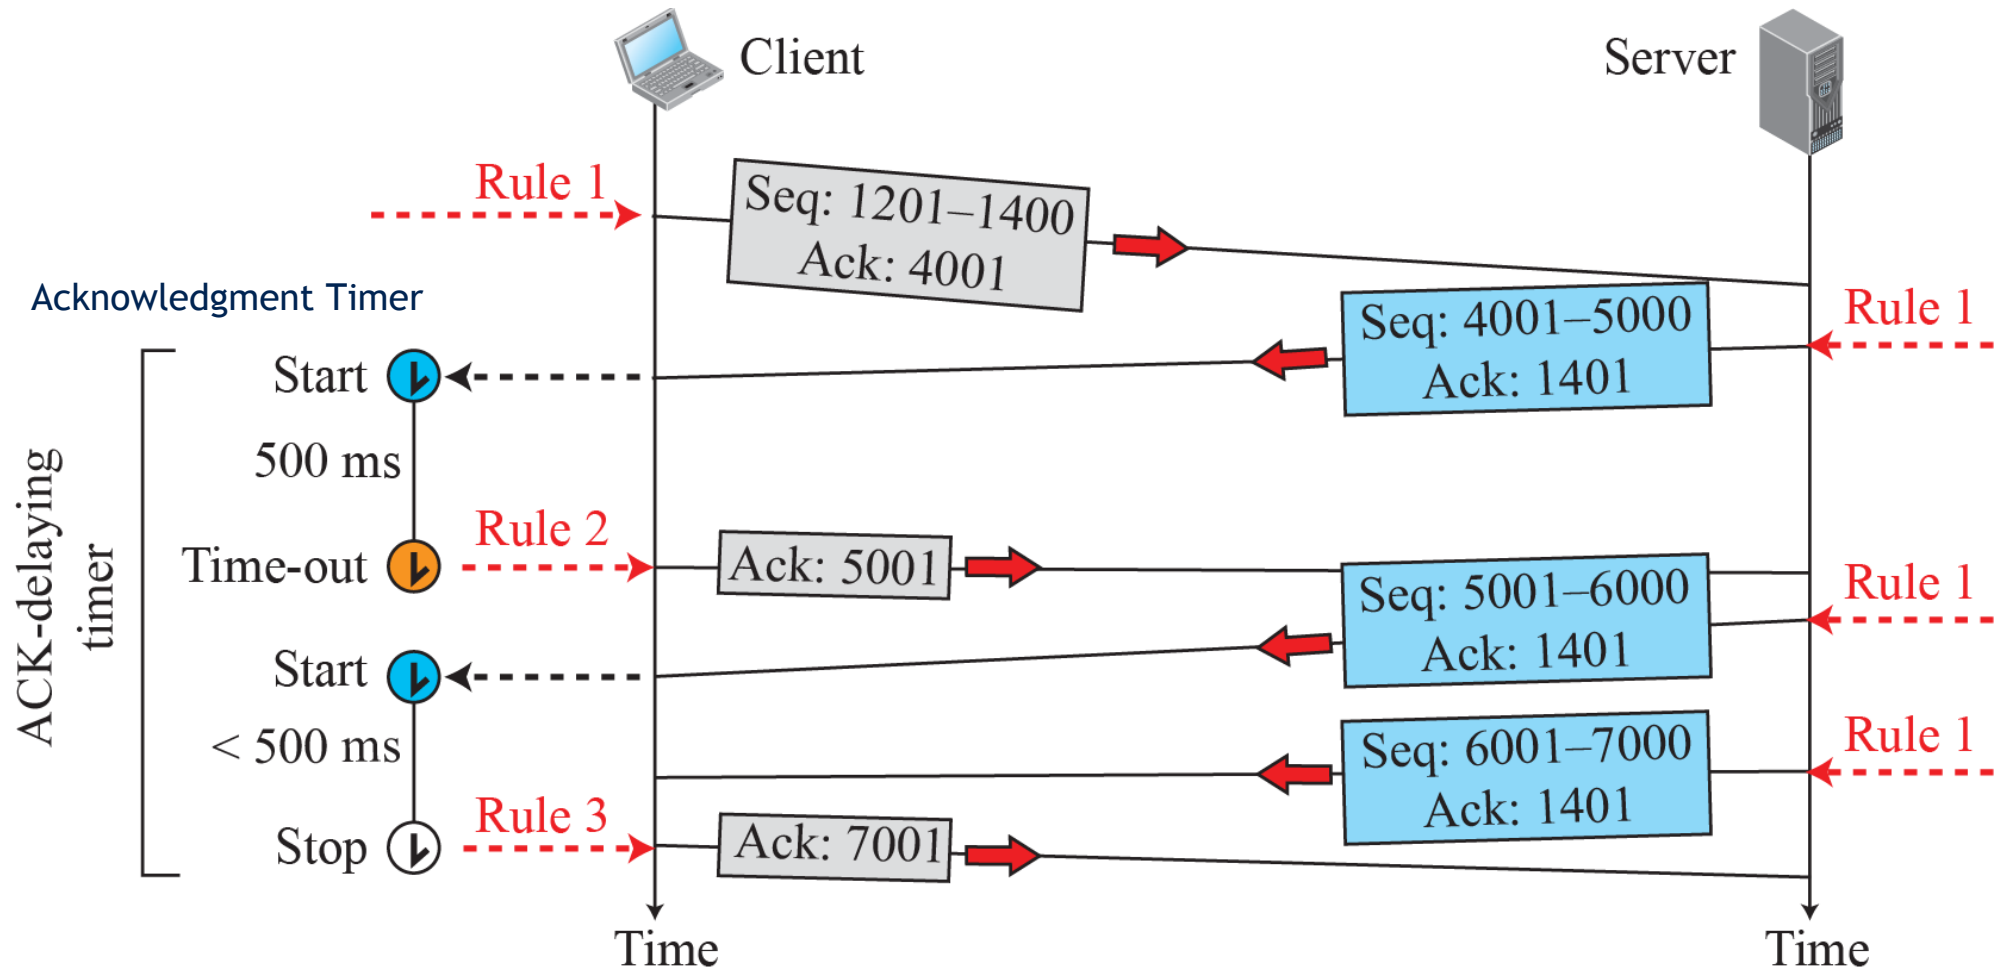
Two Army Problem

- TCP "solution" to this problem uses timeouts
- Usually works in practice

TCP Error Control

- To handle lost or damaged segments, timers are set by both sides to determine whether or not to retransmit data or ACKs
- Rule 1: When there is data to be sent, the sender must include an ACK, with the acknowledgement number set to the next sequence number it expects to receive from the receiver
- Rule 2: When the recipient has received a valid segment from the sender, and the recipient has no data to send in return, it must set a timeout after which an ACK must be returned to the sender, with the acknowledgement number set to the next expected sequence number
- Rule 3: When a second unacknowledged segment arrives, the recipient must return an ACK to the sender, with the acknowledgement number set to the next expected sequence number

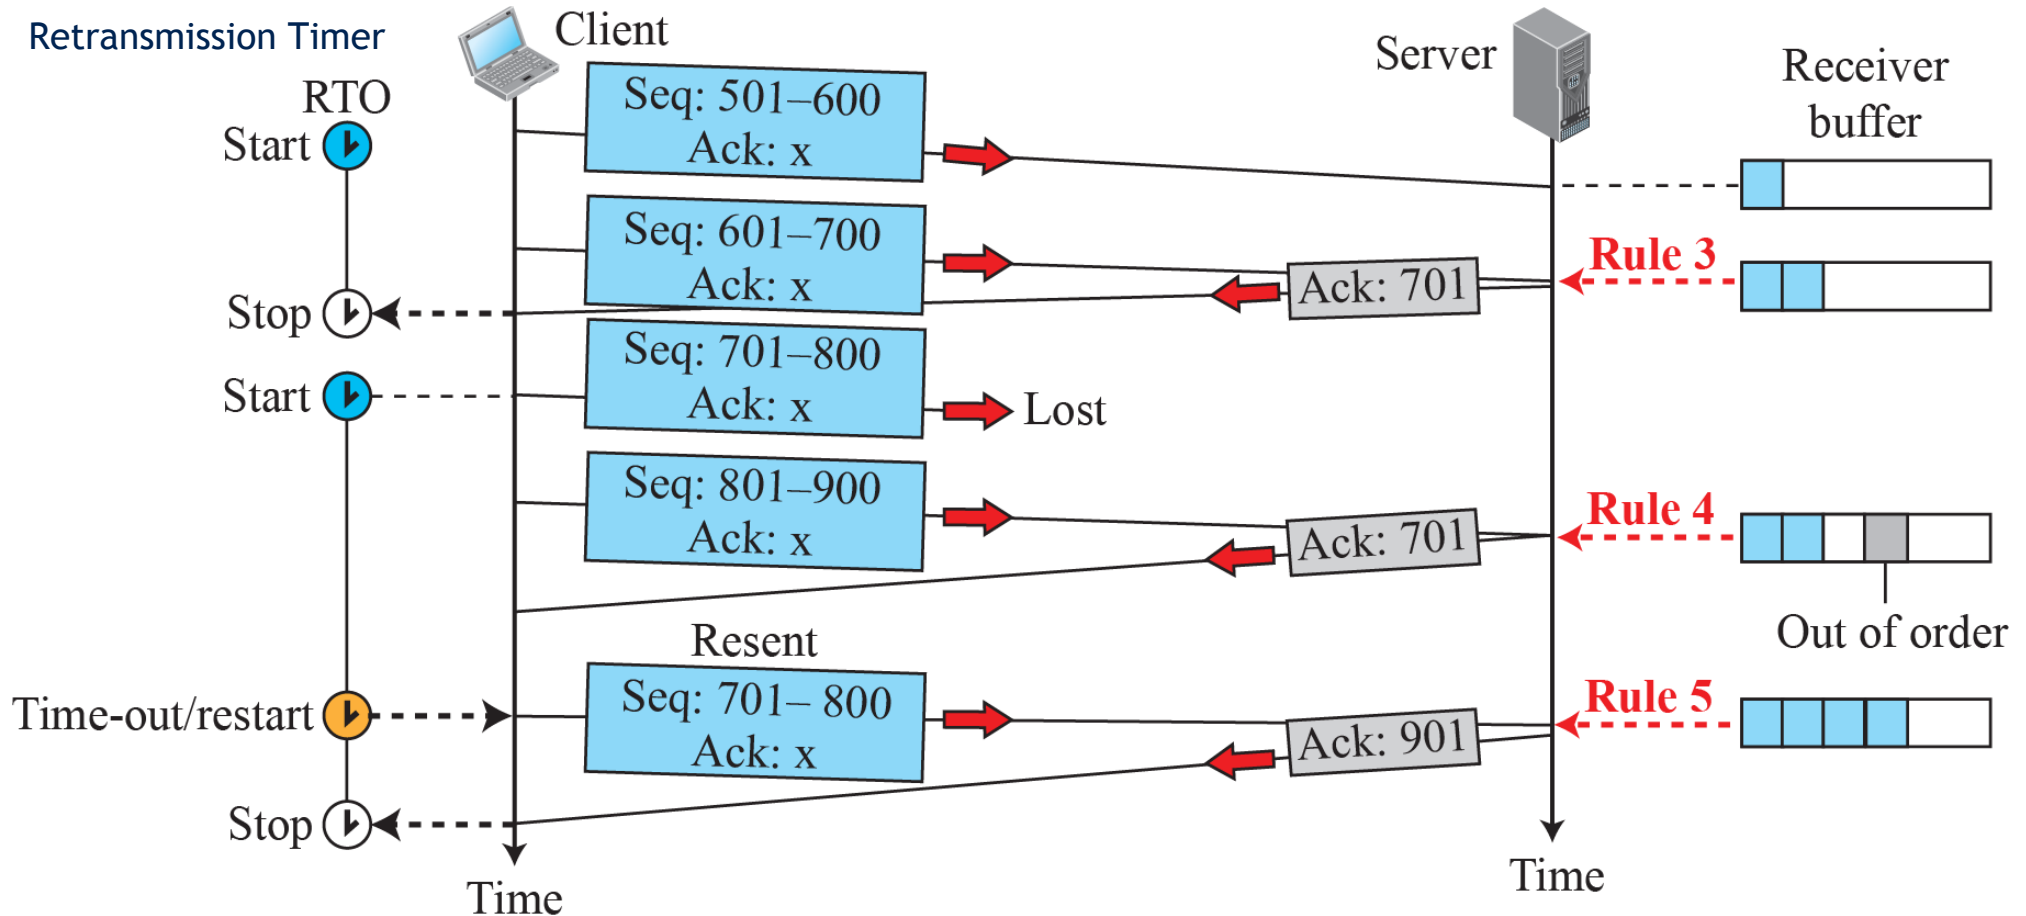
TCP Error Control



TCP Error Control

- Rule 4: When an out-of-order sequence number arrives, the segment is buffered and an ACK is returned to the sender, with the acknowledgement number set to the next expected sequence number
- Rule 5: When a missing segment arrives, the recipient returns an ACK to the sender, with the acknowledgement number set to the next expected sequence number

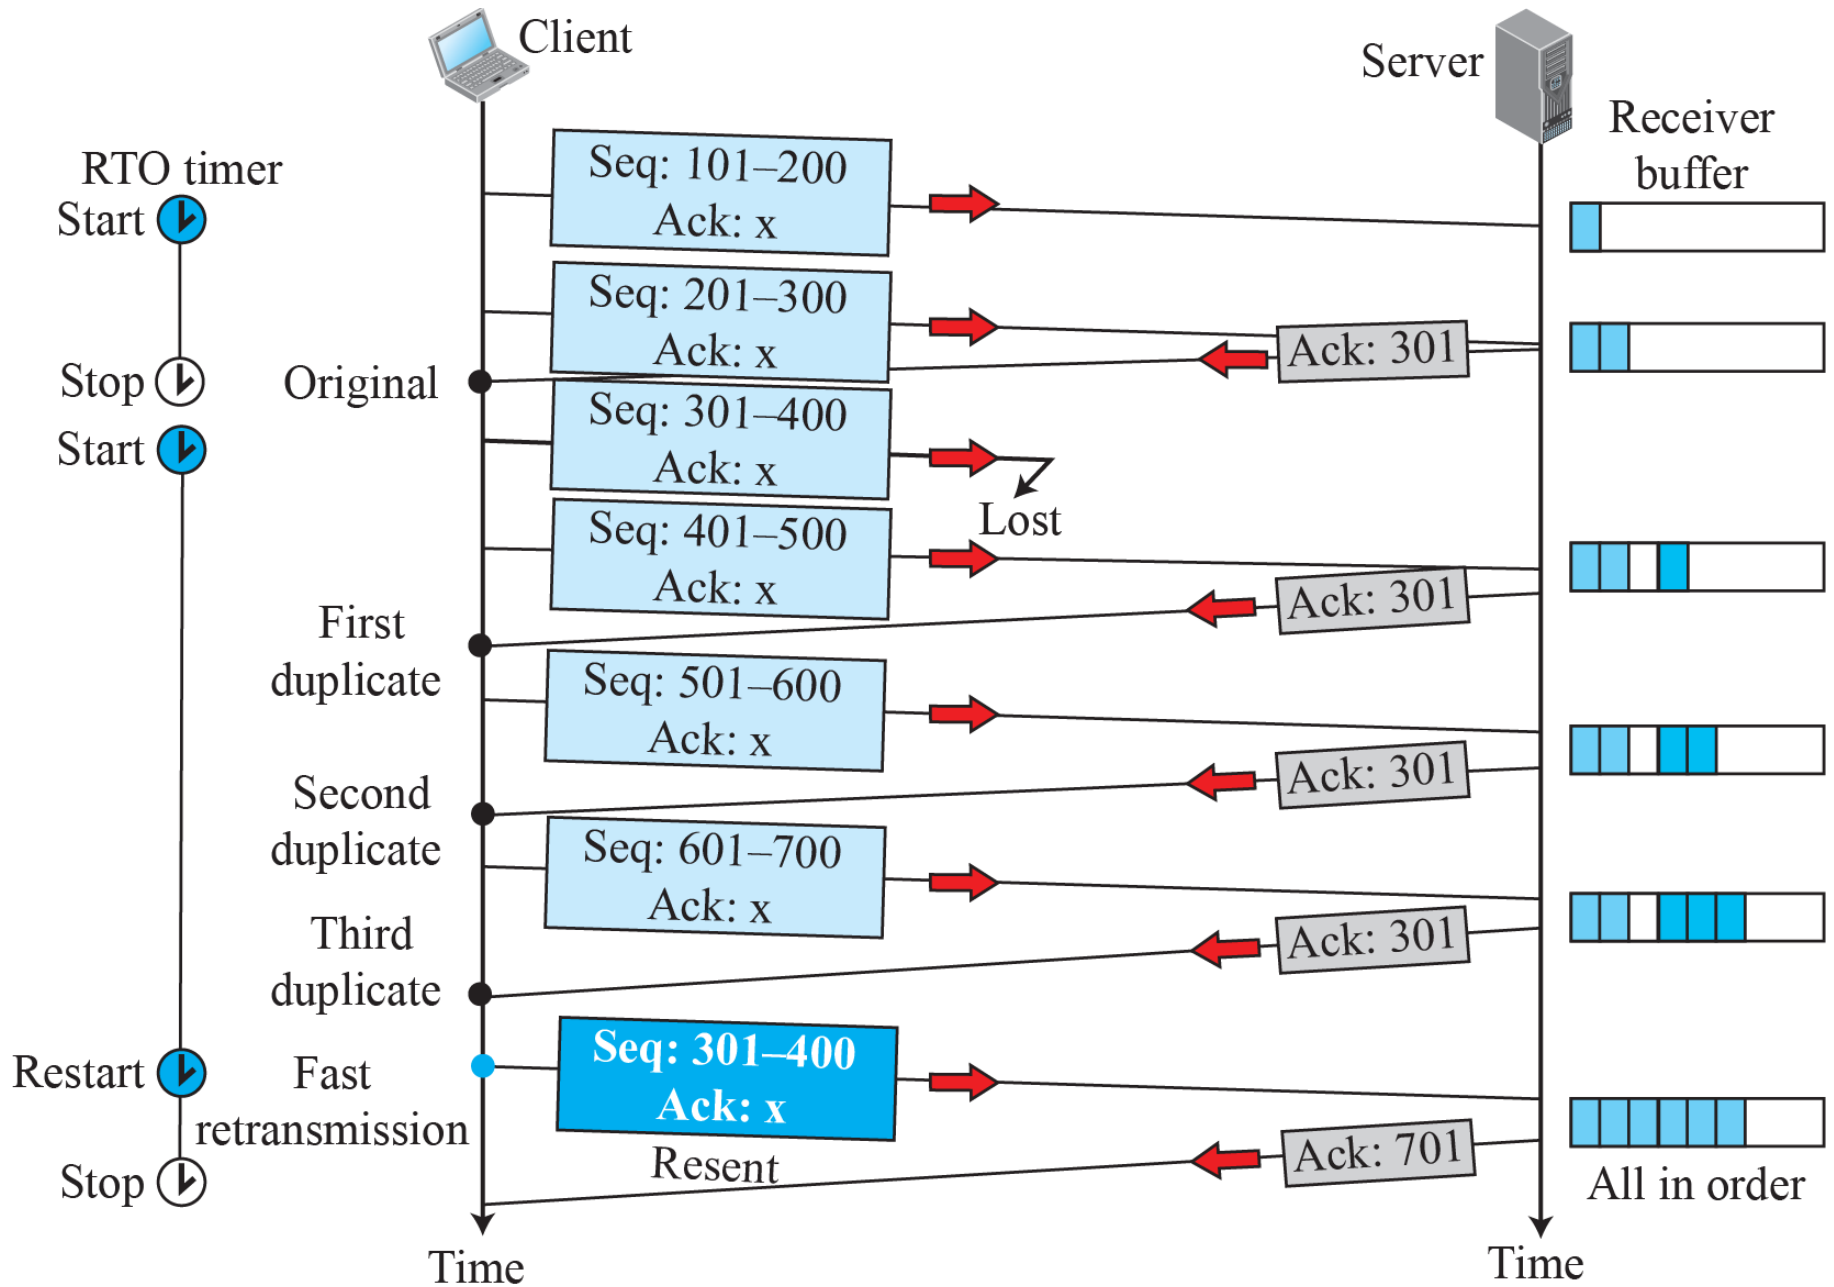
TCP Error Control



TCP Error Control

- Fast Retransmission Rule: If a 3rd duplicate ACK arrives, the subsequent segment is deemed lost, its retransmit timer is restarted, and the segment is retransmitted

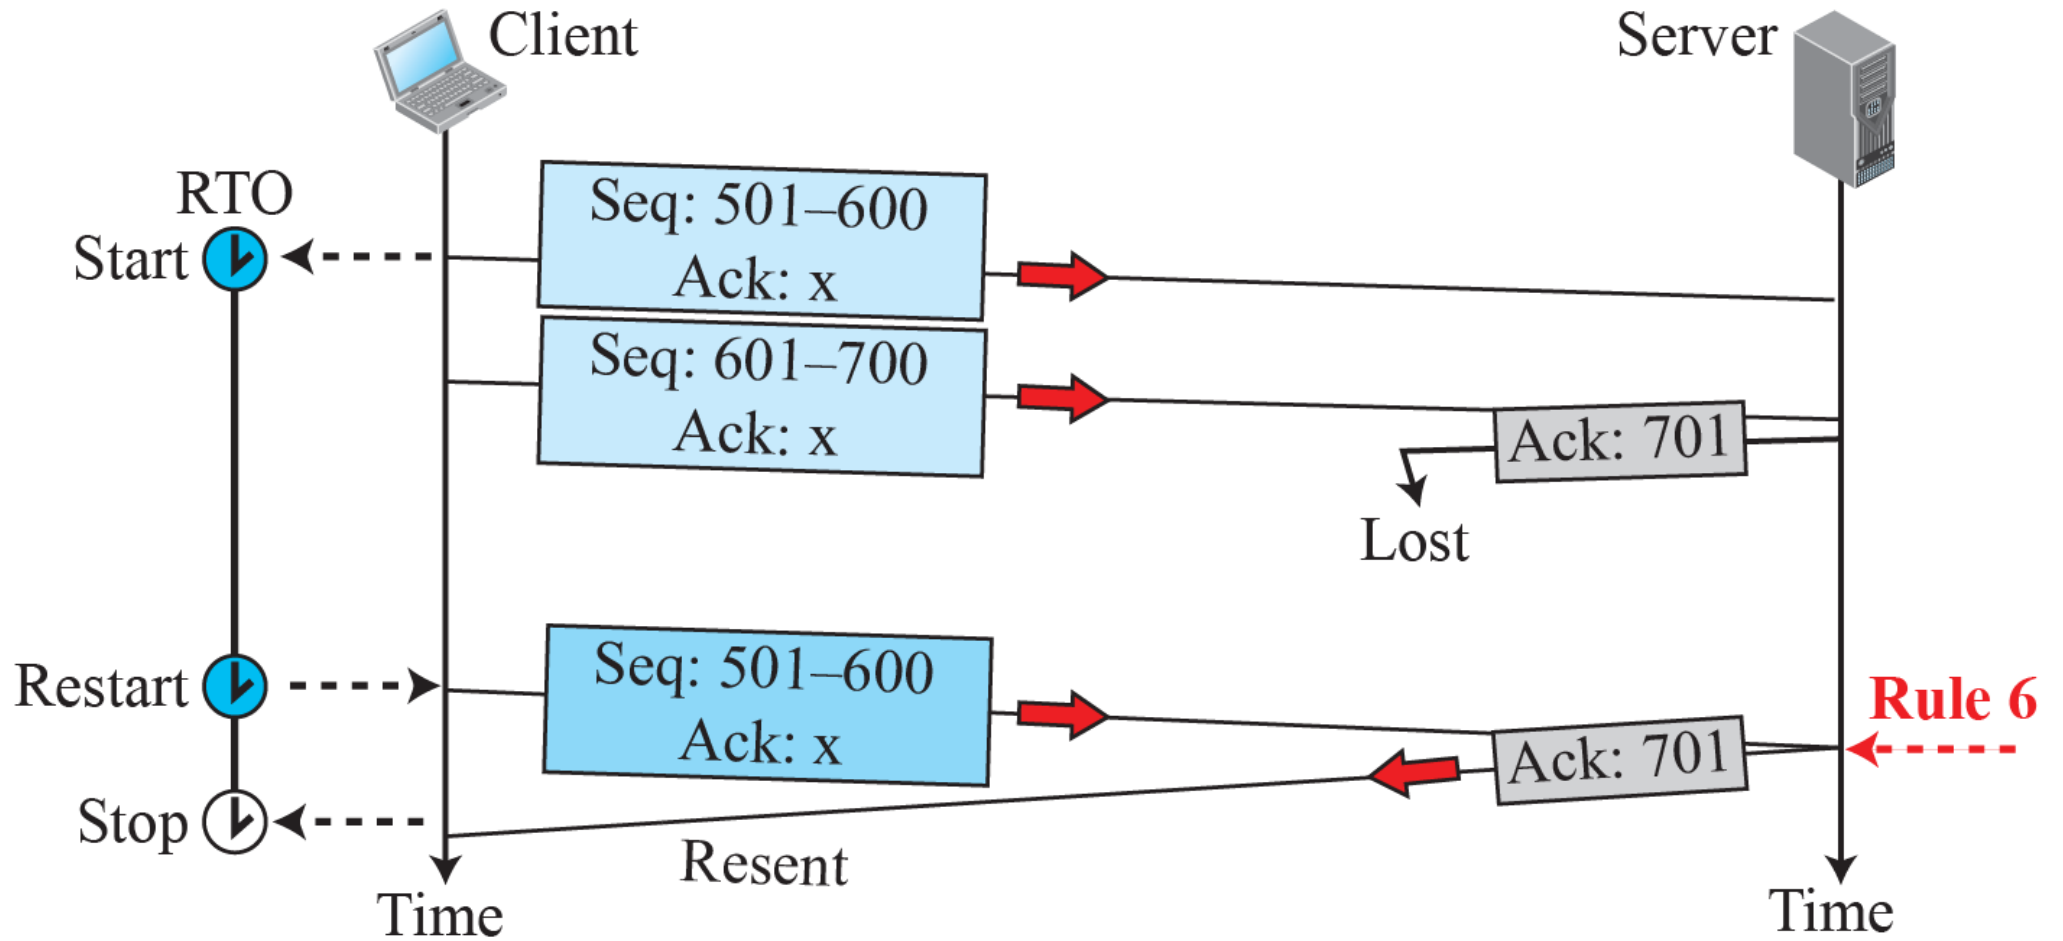
Fast Retransmission



TCP Error Control

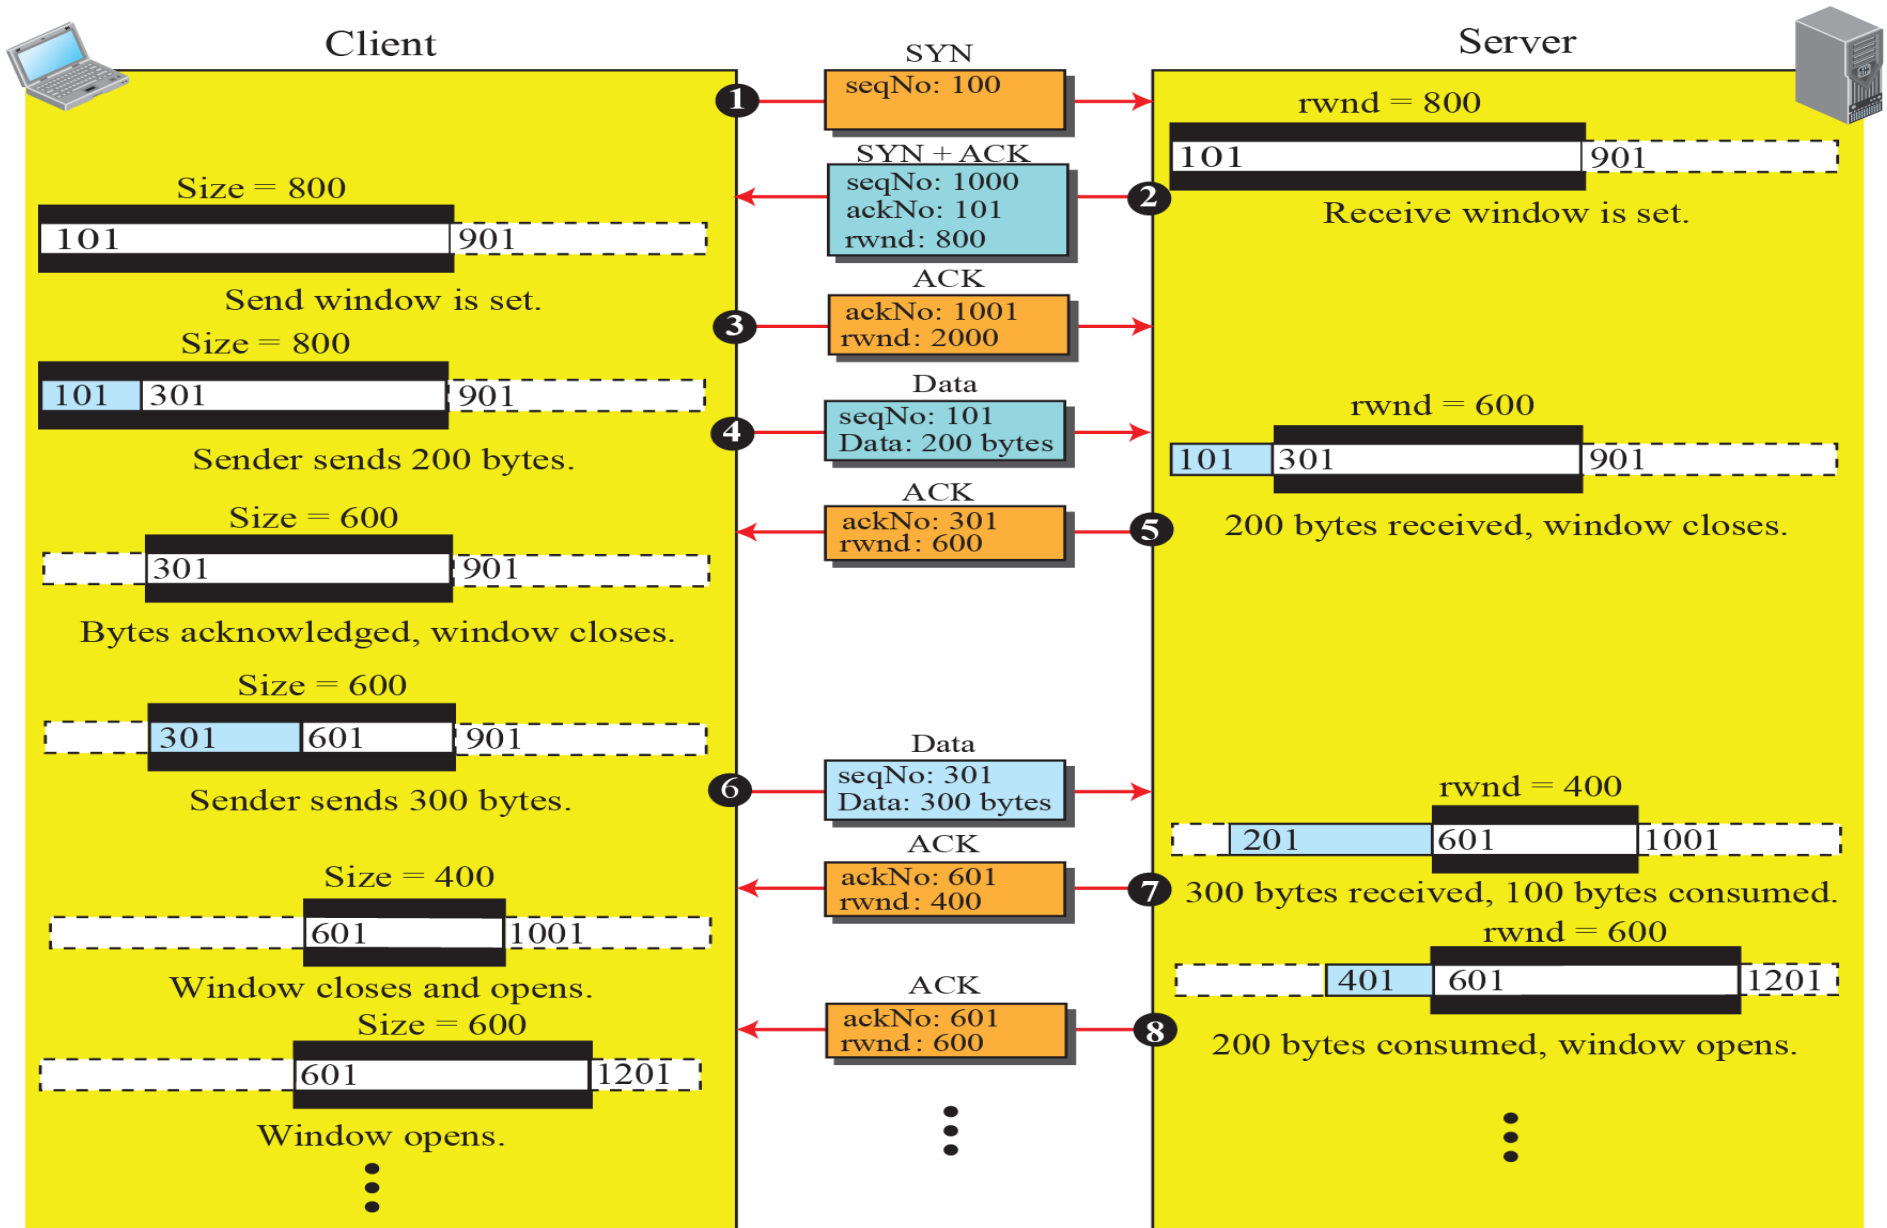
- Rule 6: When a duplicate segment is received, it is discarded, but the recipient returns an ACK to the sender, with the acknowledgement number set to the next expected sequence number

Lost ACK



TCP Flow Control

Note: We assume only unidirectional communication from client to server. Therefore, only one window at each side is shown.



Exercises

- If a valid TCP segment is received from a sender, the recipient must _____ and _____, assuming the recipient has no valid data to send

Exercises

- If a duplicate TCP segment is received from the sender, the recipient must _____, _____ and set the acknowledgement number to _____

Exercises

- Assuming the receive buffer already contains bytes 1000 - 1200 and a TCP segment arrives with sequence number 1400 and length 100, the recipient must _____ and set the acknowledgement number to _____ and save bytes _____

Key Skills

- Explain TCP connection establishment, data transmission, error control, flow control, and connection termination